VeriFact — Demo Video Script & Shot List

Superjoin VIT 2026 Engineering Intern assignment submission

Target length: 5–7 minutes

Format: Screen recording, narrated (voiceover or live commentary)

Before recording: Server running (`uvicorn app.main:app --reload`), browser at `http://127.0.0.1:8000`

> **Golden rule for this recording:** everything shown must be a real, live result from the running system — no slides pretending to be the app, no pre-scripted "expected" output. If a case resolves differently than described here on the day of recording (e.g. re-ingested facts came out slightly differently), narrate the ACTUAL result honestly rather than re-doing the take to match this script. That honesty is itself part of what this project is trying to demonstrate.

Scene 1 — Hook & the problem (0:00–0:30)

	SHOW	SAY
0:00	Two documents side by side (or the README's opening example): an annual report stating revenue as ₹81,415.38 million, an earnings deck stating ₹8,142 crore.	"Two documents from the same company. One says revenue was ₹81,415.38 million. The other says ₹8,142 crore. Are these the same number, or a real discrepancy? Nobody currently checks — VeriFact does."

Scene 2 — Architecture in 30 seconds (0:30–1:15)

	SHOW	SAY
0:30	README's pipeline diagram, or draw it live: PDF → extraction → evidence → candidates → semantic match → deterministic checks → adjudication → graph coherence.	"VeriFact reads PDFs, extracts facts tied to verbatim quotes, and figures out whether facts across documents agree, disagree, or disagree for a good reason. The core idea: the LLM proposes — deciding things that need language understanding, like whether 'net worth' and 'total equity' mean the same thing. Deterministic code verifies everything that has a computable answer, and it does not just ask nicely — it checks."

Scene 3 — Live upload (1:15–1:45)

	SHOW	SAY
1:15	Upload tab → drag in a Delhivery PDF (or click an already-queued one) → Documents list showing "processing" with the live progress bar (extracting → comparing).	"Uploading a real PDF. No document-specific code runs here — the same extraction prompt, the same pipeline, regardless of what's inside."

> **Tip:** if a fresh upload would take too long on camera, cut to an already-completed document and say so plainly on camera ("this one already finished processing, so I'll jump to a completed document") rather than implying it

happened instantly.

Scene 4 — Case 1: Corroboration (1:45–2:30)

Relationships tab → filter to a corroborates relationship (the ₹81,415.38M / ₹8,142 Cr pair)

1:45	**SHOW** The relationship card: both facts with their quotes, the green "Values agree once normalized" strip, the `corroborates` badge, and the small "code confirmed the model's answer" text.	**SAY** "Same number, two units. ₹81,415.38 million and ₹8,142 crore reduce to the same base value in code — not by asking the model to do the arithmetic, by computing it. And notice this line: 'code confirmed the model's answer.' The model proposed corroborates too, but the label you're seeing is backed by a computed 0.006% agreement, not just the model's word."

Scene 5 — The signature feature: LLM vs. system disagreement (2:30–3:45)

Relationships tab → filter to "Related, not comparable" or "Uncertain" → find a card with the yellow "system overrode its own model" banner

2:30	**SHOW** A relationship card where the yellow banner reads "The system overrode its own model here" — the model's original proposal shown in bold, the deterministic reason underneath.	**SAY** "This is the part I actually want to highlight. Here, the model proposed one label — but a deterministic check caught something the model didn't: [read the actual disagreement_reason on screen]. This is stored and shown, never hidden. Every relationship carries decision_source: whether a computed check confirmed the model, overrode it, or — for genuinely qualitative facts with no number to check — deferred to the model's own reading, honestly labeled as such rather than dressed up as equally certain."
3:15	**SHOW** Knowledge Graph tab → Logic Check tab → an "Impossible triangle" card, pointing at the "strict proof" vs. "strong evidence, model-judged equality" tag.	**SAY** "The graph checks itself too. If A agrees with B, and B agrees with C, then A must agree with C — that's transitive. When a triangle violates that, at least one judgment is wrong, found with no ground truth and no extra model call. And we're precise about how strong that claim is: strict proof when the equality is backed by numbers, strong evidence — not literally a proof — when it rests on the model's own qualitative reading."

Scene 6 — Case 3: Reconciled by context (3:45–4:20)

A standalone vs. consolidated revenue pair, relation_type = reconciled		
3:45	**SHOW** The reconciled card, the "Reconciled by: scope (consolidated vs standalone)" line.	**SAY** "These two values genuinely disagree — about 7-8%. But one is standalone, one is consolidated. That's a real, expected difference, determined in code from the scope field, not guessed by the model. Reconciled, not a false alarm."

Scene 7 — Case 4: Failures, found and surfaced (4:20–5:15)

Extraction Issues tab (pick a document from the sidebar filter) → a `value_not_evidenced` row, then a coherence violation		
4:20	**SHOW** An extraction issue: a quote that's genuinely on the page but doesn't support the claimed value (a chart axis label or table row-label case).	**SAY** "Not every extraction is clean, and VeriFact doesn't pretend otherwise. Here the quote IS real — but it's a chart's axis labels, not an actual figure. The system marks this fact as evidenced-but-unvalidated rather than confidently wrong."
4:50	**SHOW** Use the new per-document filter on the Extraction Issues / Logic Check tabs to scope to one document.	**SAY** "Every view — Facts, Relationships, Priority, Timelines, Logic Check, Extraction Issues — can be scoped to a single document, so you can audit one file's results in isolation."

Scene 8 — Generalization: an unseen document (5:15–6:15)

Sidebar → an IMF or RBI document (not Delhivery) → its Facts and a Timeline built from it		
5:15	**SHOW** The IMF India Article IV or RBI Annual Report document in the sidebar — genuinely different domain, different table layout.	**SAY** "Delhivery is a demo corpus, not a dependency. This is an IMF macroeconomic report and an RBI annual report — documents the code was never written with in mind. Same pipeline, zero document-specific code."
5:45	**SHOW** Timelines tab, filtered to the RBI document → the Foodgrains Production or inflation-rate trend line.	**SAY** "And it still works: here's a real trend the system found on its own — foodgrains production across two fiscal years — built from the exact same relationship-chaining logic that was tuned entirely on Delhivery's revenue figures."

Scene 9 — Wrap-up (6:15–6:45)

	SHOW	SAY
6:15	Overview dashboard — the corpus-wide stats (facts, relationships, evidence quality donut).	"668 automated tests, none of them needing a live model. Every claim in the README is a measured number, not an estimate — including where things didn't work, like the pages that timed out on this local 8B model. The full write-up, including everything found during a generalization pass on unseen documents, is in FINAL_REVIEW.md."

Total runtime target: ≈ 6:45

Pre-recording checklist

- Server started fresh (`uvicorn app.main:app --reload --port 8000`) so the UI reflects the current code.
- At least one Delhivery document and both the IMF and RBI documents already ingested and visible in the sidebar (ingestion itself is slow on a local model — don't record dead air waiting for it unless that pause IS the point being made about performance).
- Confirm a real `disagreement: true` relationship exists to show in Scene 5 — filter Relationships by "Related, not comparable" or "Uncertain" beforehand and note which document/fact pair to click to.
- Confirm a real coherence violation is visible on the Logic Check tab.
- Browser zoom / window size set so text is legible in the recording (100–110% zoom, 1440px+ width recommended).
- Close any unrelated browser tabs/notifications before starting the capture.

Generated as a planning aid for the recorded submission video — not a transcript to read verbatim. Adapt timings and phrasing naturally; the important constraint is that every screen shown is a real, live result from the running system, matching FINAL_REVIEW.md and README.md.